

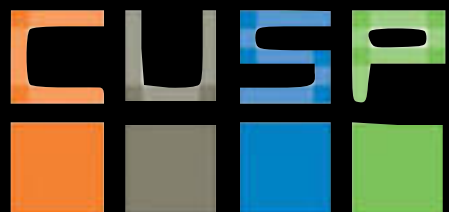
# Urban Informatics

Fall 2015

dr. federica bianco [fb55@nyu.edu](mailto:fb55@nyu.edu)

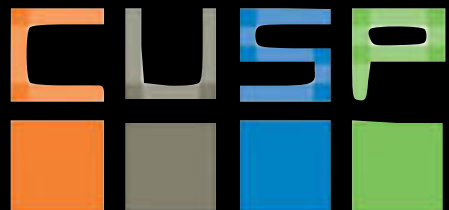


@fedhere



XII categorical distances  
model diagnostics, OSMnx

Last Class!!!!



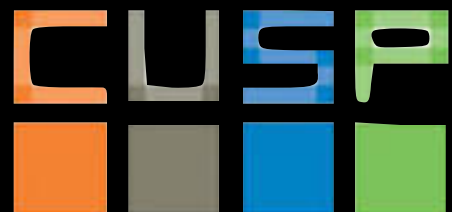
XII categorical distances  
model diagnostics, OSMnx

## Topics covered:

- Good practices with data: falsifiability, reproducibility
- Basic data retrieving and munging: APIs, Data formats
- SQL
- Basic statistics: distributions and their moments
- Hypothesis testing:  $p$ -value, statistical significance
- Statistical and Systematic errors
- Goodness of fit tests
- Visualizations
- Geospatial analysis
- Likelihood
- OLS
- Topics in (time) series analysis
- Clusters
- Decision and regression trees (CART)

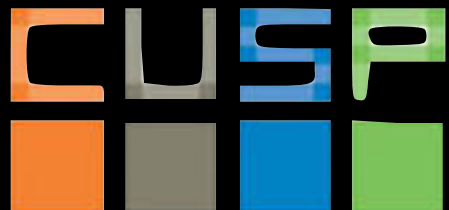
## Today:

- categorical and mixed clustering
- model diagnostics (ROC, AUC)
- OSMnx
- tips on efficient coding



XII categorical distances  
model diagnostics, OSMnx

# Distances

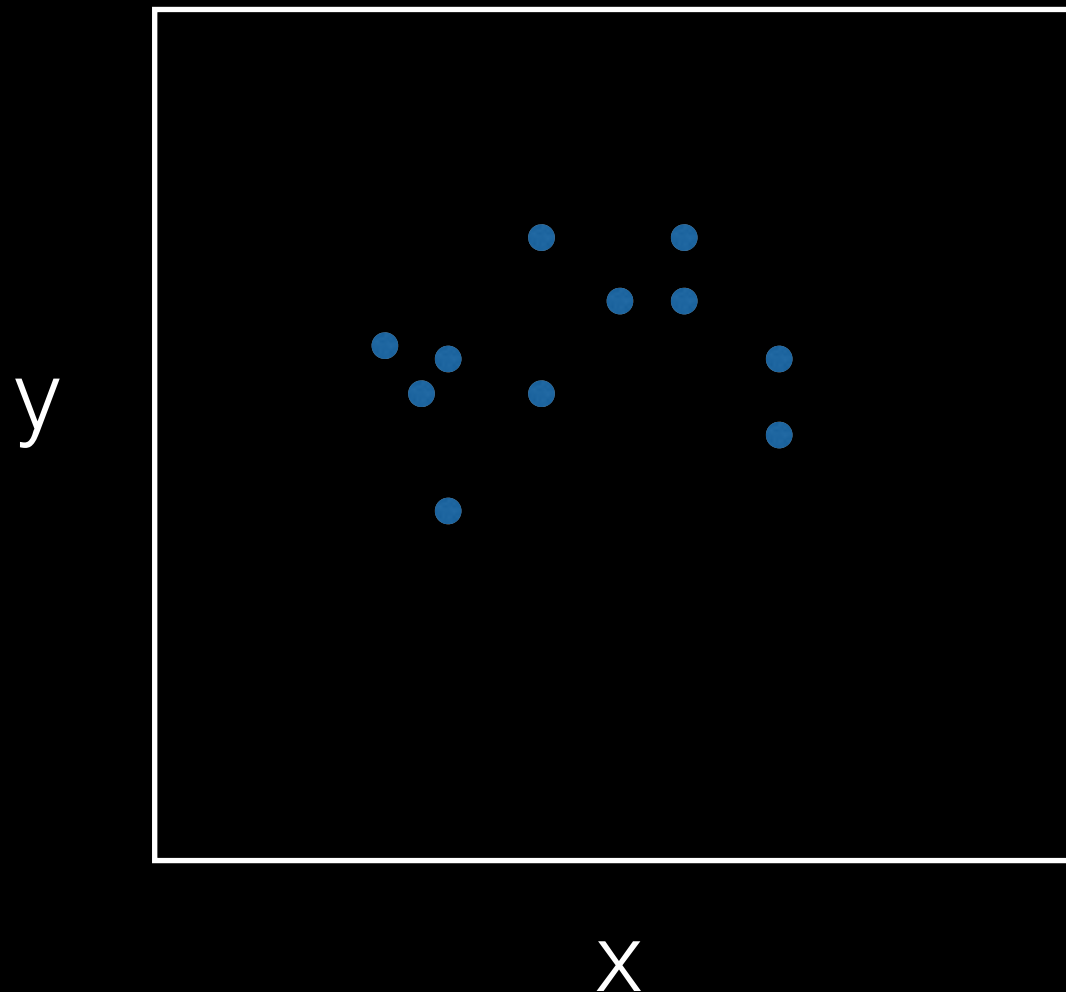


XII categorical distances  
model diagnostics, OSMnx

## Partitioning methods: clustering

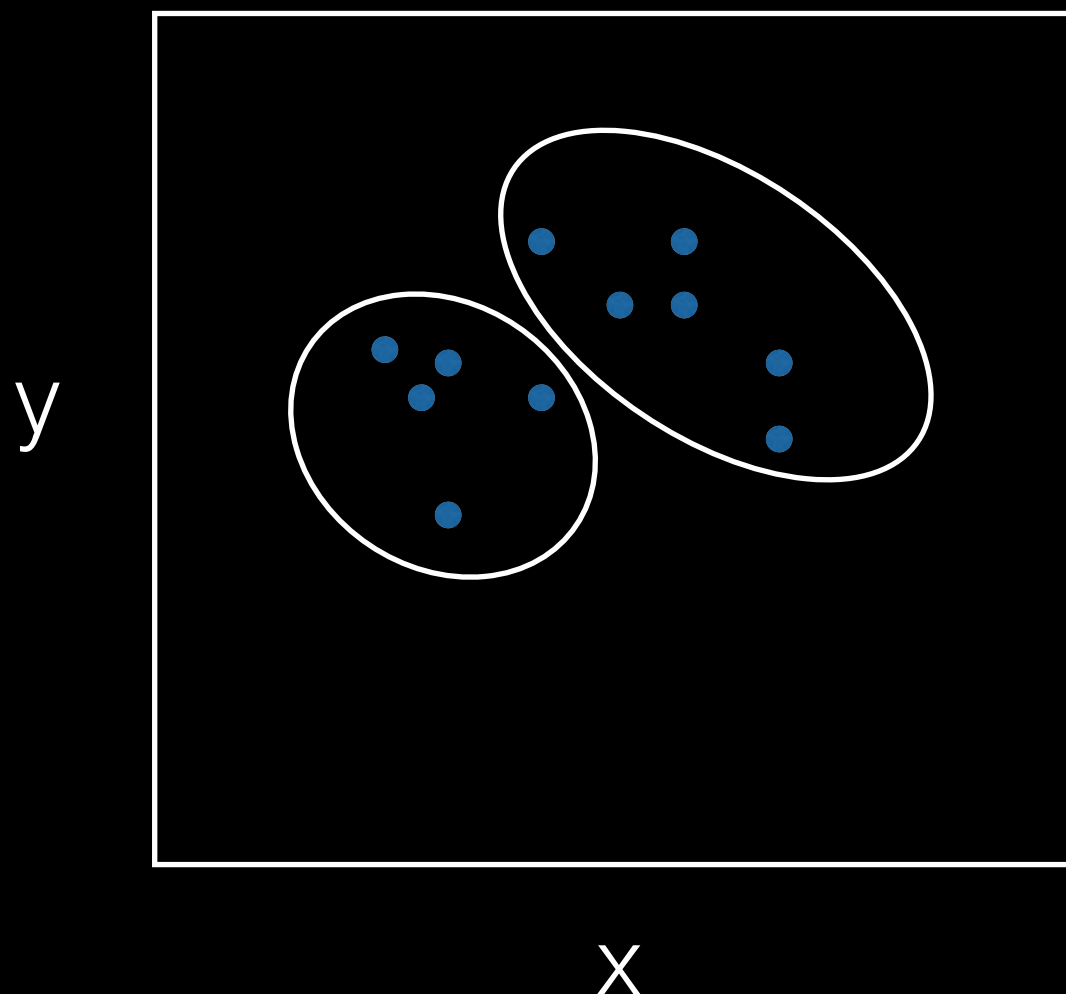
goal is to partition the space so that the observed variables are separate in maximally homogeneous groups

***observed:***  
 $(x, y)$



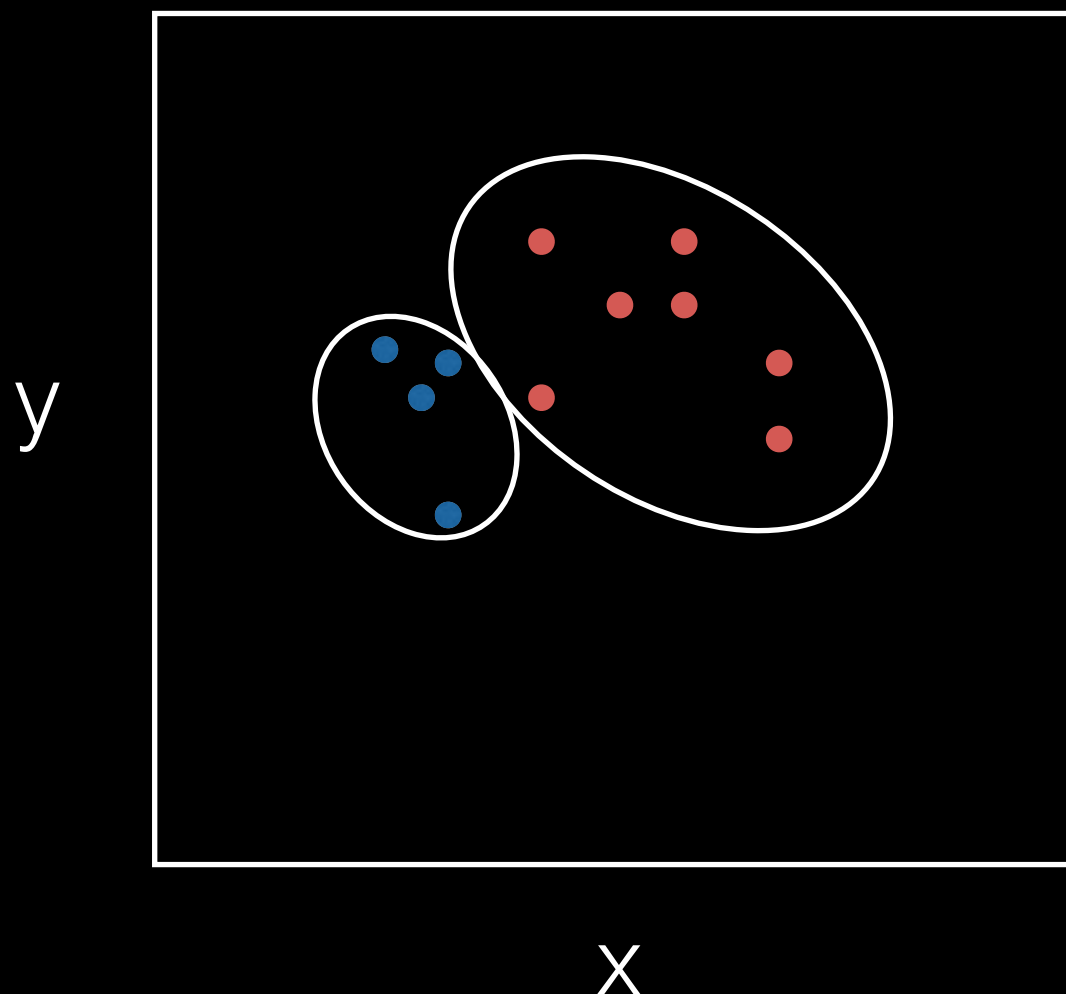
# Partitioning methods: clustering

***observed:***  
 $(x, y)$

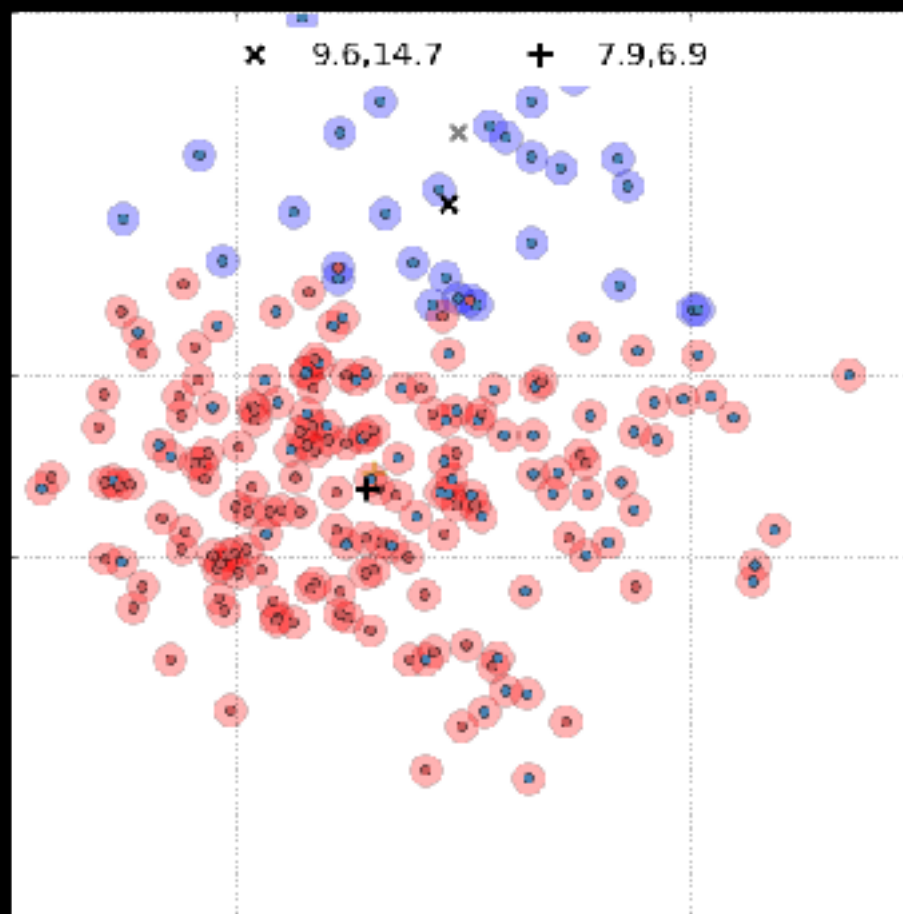


# Partitioning methods: clustering

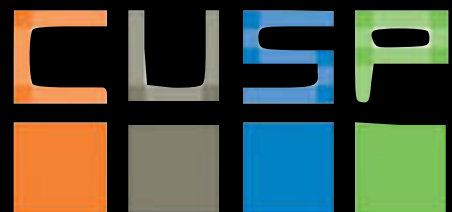
***observed:***  
(x, y, color)



# Crisp (or hard) clustering - K-means



You guess the centers and assign points to clusters  
based on a predefined distance metric

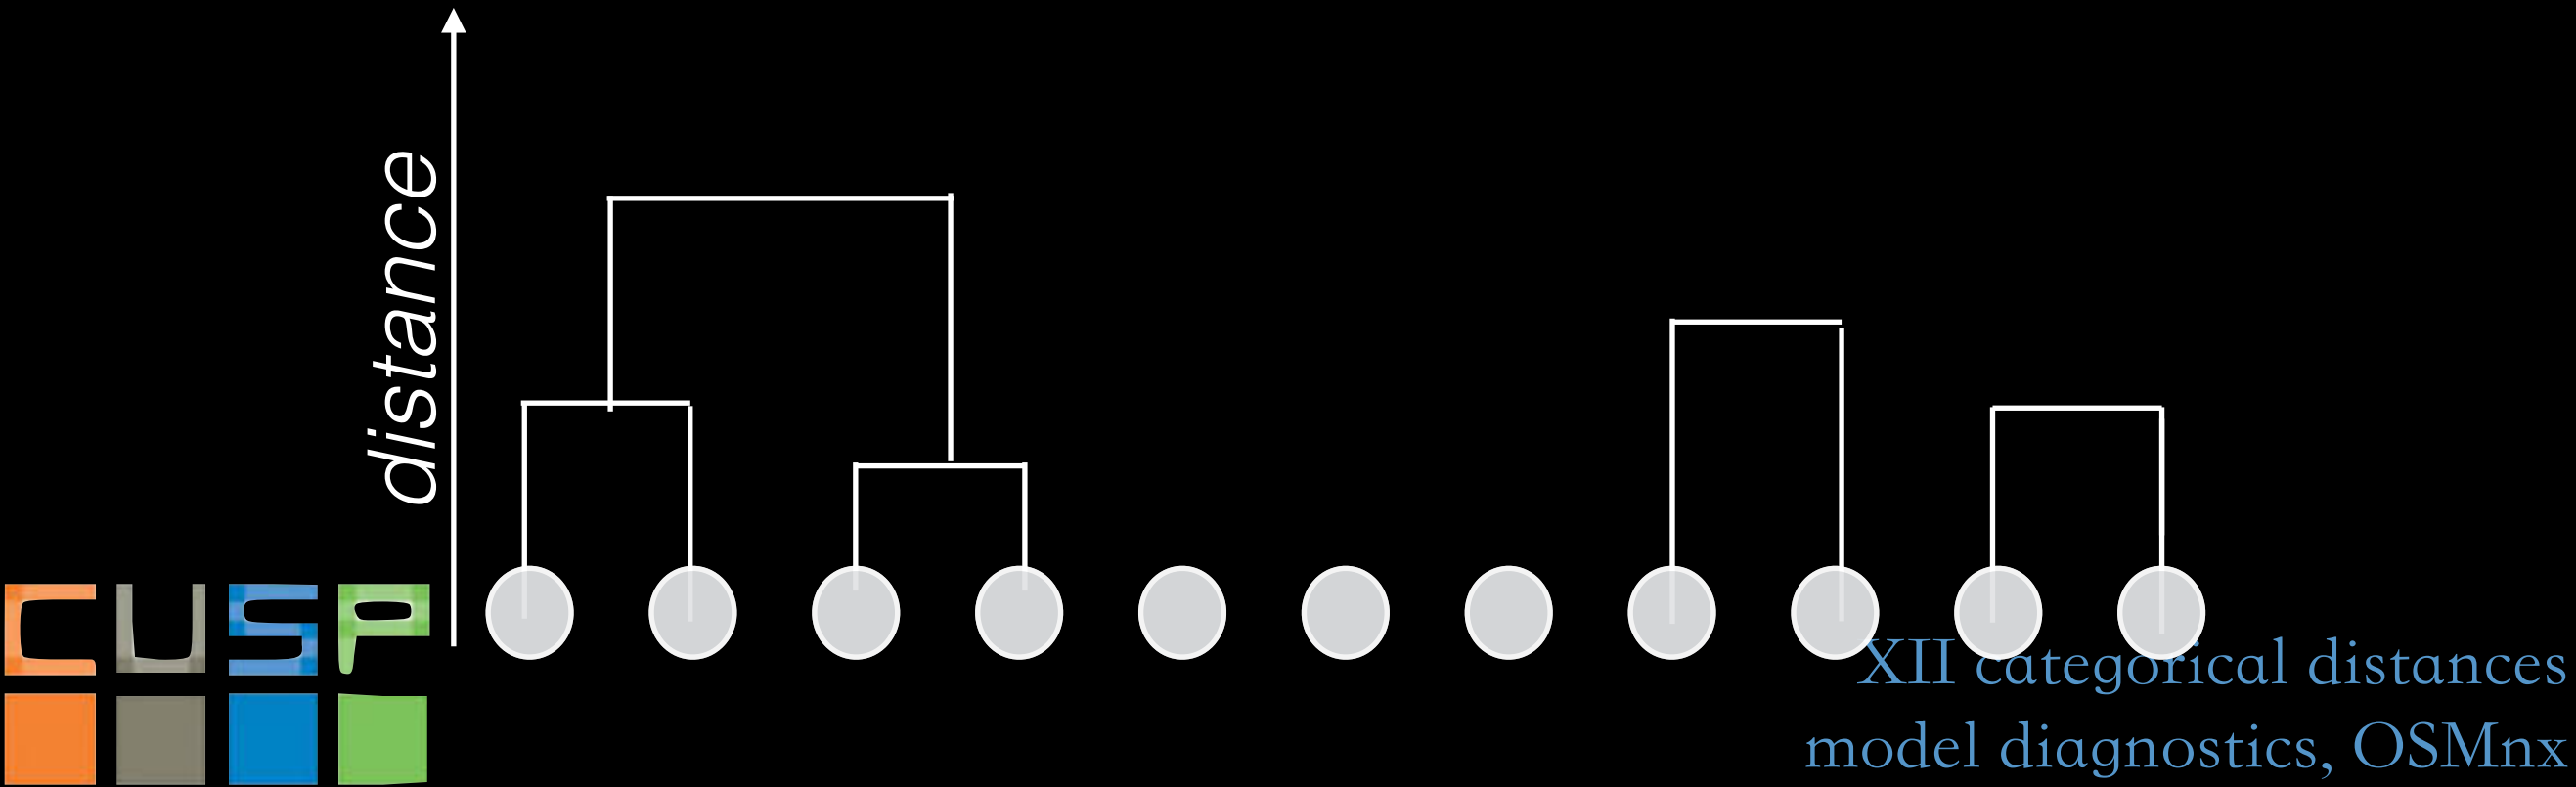
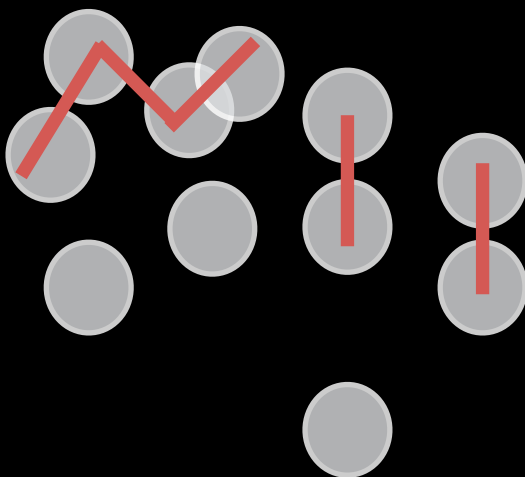


XII categorical distances  
model diagnostics, OSMnx



# hierarchical clustering

*agglomerative*  
*bottom-up*



## Summary and Key concepts

***clustering is easy, but interpreting results is tricky***

Distance metrics:

- Eucledian and other Minchowski metrics

- geospacial distances

- metrics for non continuous data

Partitioning methods: inexpensive, typically non deterministic

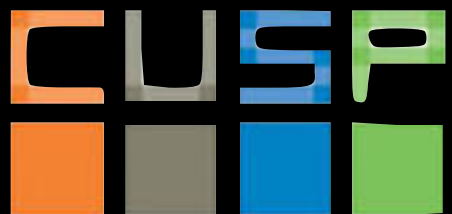
- Hard methods: *K-means, K-medoids*

- Soft (or fuzzy) methods: (i.e. probabilistic approach)

  - Expectation Maximization Mixture models*

Hierarchical methods:

- divisive vs agglomerative, dendrograms



XII categorical distances  
model diagnostics, OSMnx

## Distance Metrics Continuous variables

### Minkowski family of distances

$$D(i,j) = \sqrt[p]{\sum_{k=1}^N |x_{ik} - x_{jk}|^p}$$

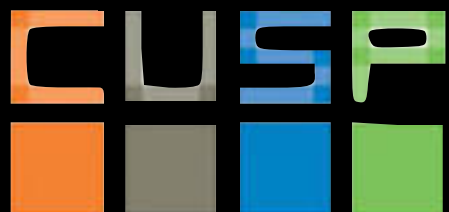
N features (dimensions)



Great Circle distances:  $\phi_i, \lambda_i, \phi_j, \lambda_j$

geographical latitude and longitude

$$D(i,j) = R \arccos(\sin \phi_i \cdot \sin \phi_j + \cos \phi_i \cdot \cos \phi_j \cdot \cos(\Delta \lambda))$$



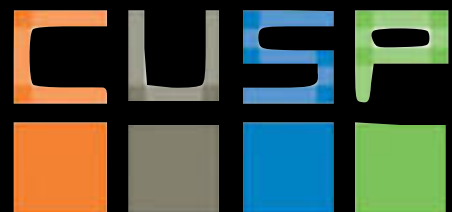
XII categorical distances  
model diagnostics, OSMnx

## Distance Metrics

## Binary variables

contingency table

	1	0	<i>sum</i>
1	<i>a</i>	<i>b</i>	<i>a+b</i>
0	<i>c</i>	<i>d</i>	<i>c+d</i>
<i>sum</i>	<i>a+c</i>	<i>b+d</i>	<i>p</i>



## Distance Metrics

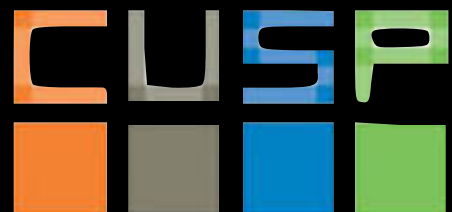
## Binary variables

contingency table

	1	0	<i>sum</i>
1	<i>a</i>	<i>b</i>	<i>a+b</i>
0	<i>c</i>	<i>d</i>	<i>c+d</i>
<i>sum</i>	<i>a+c</i>	<i>b+d</i>	<i>p</i>

e.g.: subway station

w ESCALATOR Y/N  
w ELEVATOR Y/N



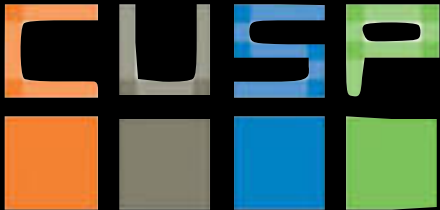
XII categorical distances  
model diagnostics, OSMnx

**Distance Metrics**    **Binary variables**  
contingency table

	1	0	<i>sum</i>
1	<i>a</i>	<i>b</i>	<i>a+b</i>
0	<i>c</i>	<i>d</i>	<i>c+d</i>
<i>sum</i>	<i>a+c</i>	<i>b+d</i>	<i>p</i>

e.g.: subway station  
w ESCALATOR Y/N  
w ELEVATOR Y/N

		ELEVATOR		
		1	0	
ESCALATOR	1	7	3	
	0	106	353	



	1	0	<i>sum</i>
1	<i>a</i>	<i>b</i>	<i>a+b</i>
0	<i>c</i>	<i>d</i>	<i>c+d</i>
<i>sum</i>	<i>a+c</i>	<i>b+d</i>	<i>p</i>

Distance Metrics

Binary variables

contingency table

e.g.: subway station

w ESCALATOR Y/N

w ELEVATOR Y/N

		ELEVATOR		
		1	0	sum
ESCALATOR	1	7	3	10
	0	106	353	459
sum		113	356	469



	1	0	<i>sum</i>
1	<i>a</i>	<i>b</i>	<i>a+b</i>
0	<i>c</i>	<i>d</i>	<i>c+d</i>
<i>sum</i>	<i>a+c</i>	<i>b+d</i>	<i>p</i>

Distance Metrics

Binary variables

contingency table

e.g.: subway station

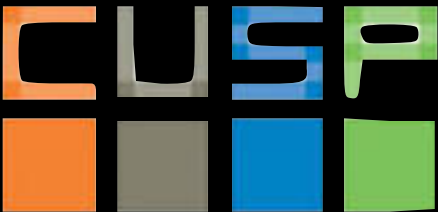
w ESCALATOR Y/N

w ELEVATOR Y/N

		ELEVATOR		
		1	0	sum
ESCALATOR	1	7	3	10
	0	106	353	459
sum		113	356	469

IF SYMMETRIC  
(same chance to appear  
i.e. roughly same total  
Y and N)

$$D_{ij} = \frac{b+c}{a+b+c+d} = \frac{109}{469} = 0.23$$





	1	0	<i>sum</i>
1	<i>a</i>	<i>b</i>	<i>a+b</i>
0	<i>c</i>	<i>d</i>	<i>c+d</i>
<i>sum</i>	<i>a+c</i>	<i>b+d</i>	<i>p</i>

Distance Metrics

Binary variables

contingency table

e.g.: subway station

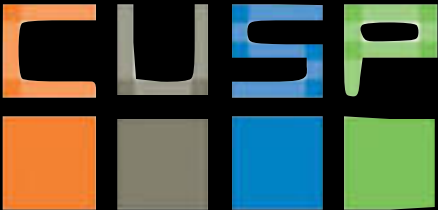
w ESCALATOR Y/N

w ELEVATOR Y/N

		ELEVATOR		
		1	0	sum
ESCALATOR	1	7	3	10
	0	106	353	459
sum		113	356	469

IF SYMMETRIC  
(same chance to appear  
i.e. roughly same total  
Y and N)

$$D_{ij} = \frac{M_{i=0j=0} + M_{i=1j=1}}{M_{00} + M_{01} + M_{10} + M_{11}} = \frac{109}{469} = 0.23$$



	1	0	<i>sum</i>
1	<i>a</i>	<i>b</i>	<i>a+b</i>
0	<i>c</i>	<i>d</i>	<i>c+d</i>
<i>sum</i>	<i>a+c</i>	<i>b+d</i>	<i>p</i>

Distance Metrics

Binary variables

contingency table

e.g.: subway station

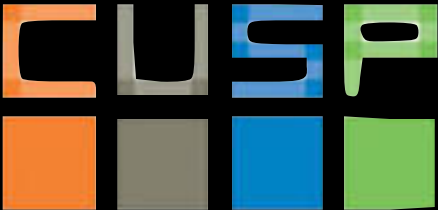
w ESCALATOR Y/N

w ELEVATOR Y/N

		ELEVATOR		
		1	0	sum
ESCALATOR	1	7	3	10
	0	106	353	459
sum		113	356	469

IF ASYMMETRIC  
(not same chance)

$$D_{ij} = \frac{b+c}{a+b+c} = \frac{109}{116} = 0.94$$



	1	0	<i>sum</i>
1	<i>a</i>	<i>b</i>	<i>a+b</i>
0	<i>c</i>	<i>d</i>	<i>c+d</i>
<i>sum</i>	<i>a+c</i>	<i>b+d</i>	<i>p</i>

Distance Metrics

Binary variables

contingency table

e.g.: subway station

w ESCALATOR Y/N

w ELEVATOR Y/N

		ELEVATOR		
		1	0	sum
ESCALATOR	1	7	3	10
	0	106	353	459
sum		113	356	469

IF ASYMMETRIC

(not same chance)

Jaccard similarity

$$J_{ij} = \frac{a}{a+b+c} = \frac{7}{116} = 0.06$$

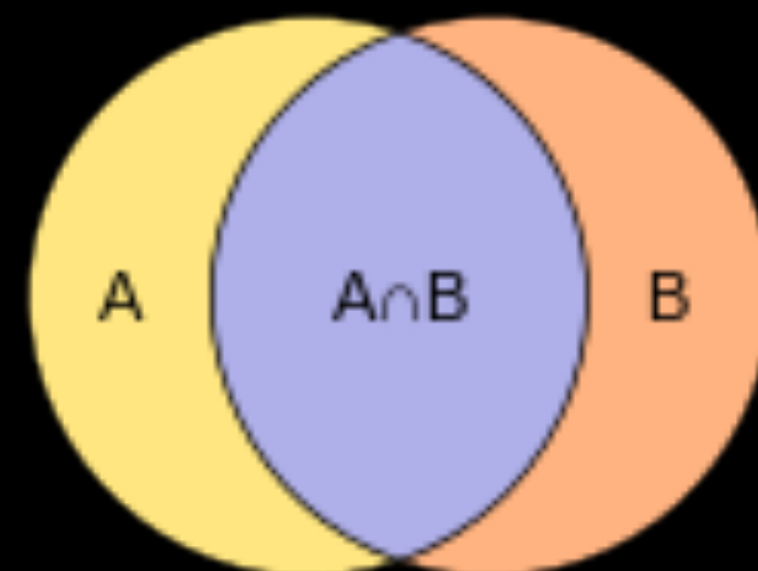


## Distance Metrics    Binary variables

Uses presence/absence data

**Jaccard similarity coefficient  $S_j$**

$$S_j = \frac{a}{a+b+c}$$



$a$  = number of items in common,

$b$  = number of items unique to the first set

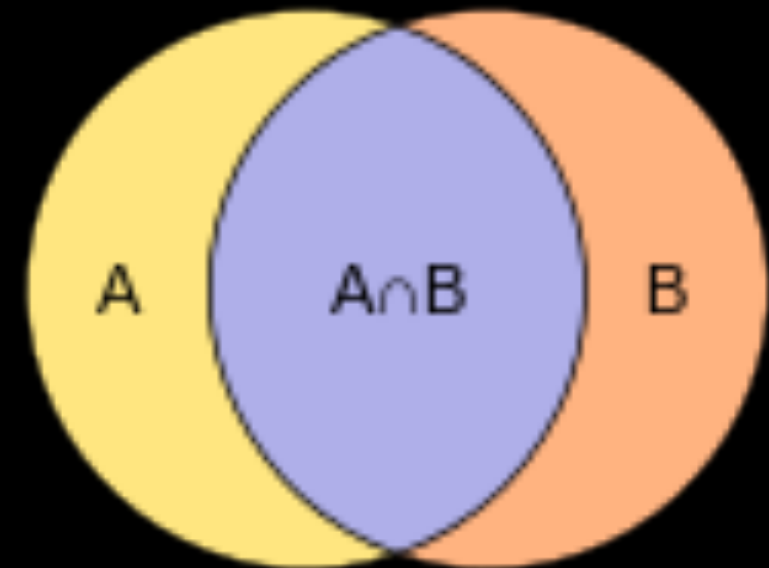
$c$  = number of items unique to the second set

## Distance Metrics    Binary variables

Uses presence/absence data

**Jaccard similarity  
coefficient  $S_j$**

$$S_j = \frac{A \cap B}{A \cup B}$$



a = number of items in common,

b = number of items unique to the first set

c = number of items unique to the second set

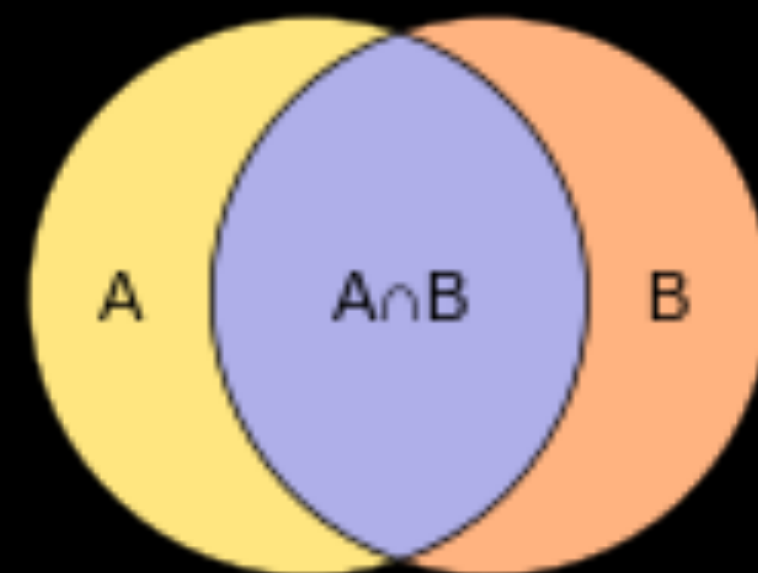
## Distance Metrics    Binary variables

Uses presence/absence data

### Jaccard distance

$$D_j = 1 - S_j$$

$$S_j = \frac{A \cap B}{A \cup B}$$



a = number of items in common,

b = number of items unique to the first set

c = number of items unique to the second set

## Distance Metrics    Categorical Variables

Uses presence/absence data in two samples (non exclusive)

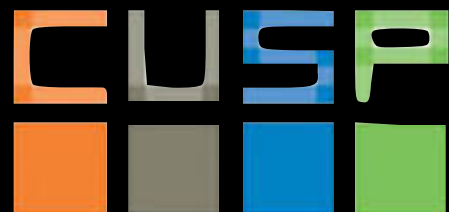
**Simple similarity coefficient**  
***Simple Matching Method***  
***SMC***

$$S_{ij} = \frac{p-m}{p}$$

$p$ : number of variables  
 $m$ : number of matches



[https://github.com/fedhere/Ulnotebooks/blob/master/cluster/categorical\\_clustering.ipynb](https://github.com/fedhere/Ulnotebooks/blob/master/cluster/categorical_clustering.ipynb)



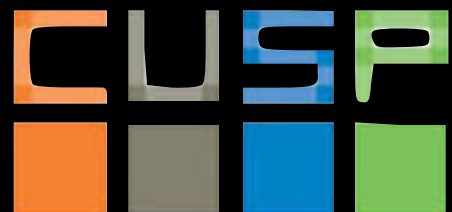
XII categorical distances  
model diagnostics, OSMnx

## Distance Metrics    Ordinal variables

Uses ranks

*map occurrences in a range 0-1*

$$r_{ij} = \{1 \dots R_N\} \rightarrow \mathbf{z}_{ij} = \frac{r_{ij} - 1}{R_N - 1}$$



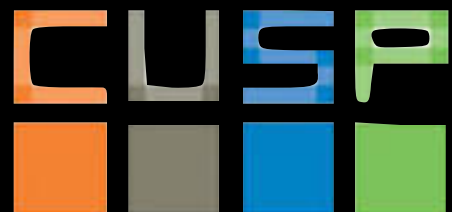


## Distance Metrics    MIXED variables

Hybrid dataset containing continuous, ordinal, categorical

***weighted distance***

$$D_w = \frac{\sum_{p=1}^p w_{ij}^{(f)} d_{ij}^{(f)}}{\sum_{p=1}^p w_{ij}^{(f)}}$$



XII categorical distances  
model diagnostics, OSMnx

## Distance Metrics    vector Variables

Uses correlation coefficient!

A time series is a vector:

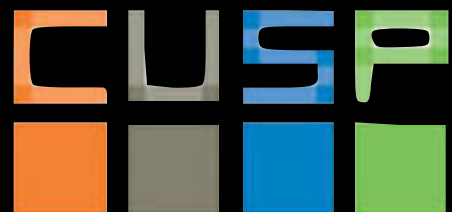
MTA rides/NYC establishments can be clustered w this distance  
clustering time series + other features requires this

### Pearson's correlation

$$\rho_{X,Y} = \frac{\text{cov}(X,Y)}{\sigma_X \sigma_Y}$$

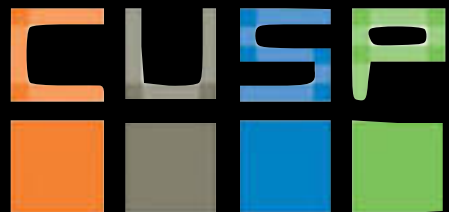
### Cosine similarity

$$\cos(\theta) = \frac{A \cdot B}{\|A\| \|B\|}$$



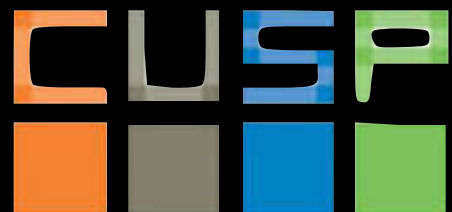
XII categorical distances  
model diagnostics, OSMnx

# Model Diagnostics



XII categorical distances  
model diagnostics, OSMnx

# Accuracy, Recall, Precision

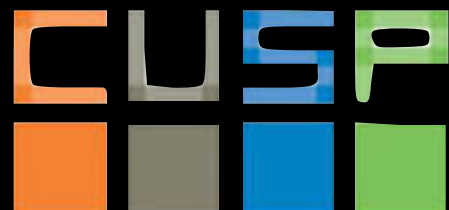


XII categorical distances  
model diagnostics, OSMnx

# Accuracy, Recall, Precision

$$LR = \frac{\text{False Negative}}{\text{True Negative}}$$

	<i>H</i> <sub>0</sub> is True	<i>H</i> <sub>0</sub> is False
<i>H</i> <sub>0</sub> is falsified	<b>Type I error False Positive</b> important message gets spammed	True Positive
<i>H</i> <sub>0</sub> is not falsified	True Negative	<b>Type II error False negative</b> Spam in your Inbox



# Accuracy, Recall, Precision

$$\text{Precision} = \frac{\text{TP}}{\text{TP} + \text{FP}}$$

$$\text{Recall} = \frac{\text{TP}}{\text{TP} + \text{FN}}$$

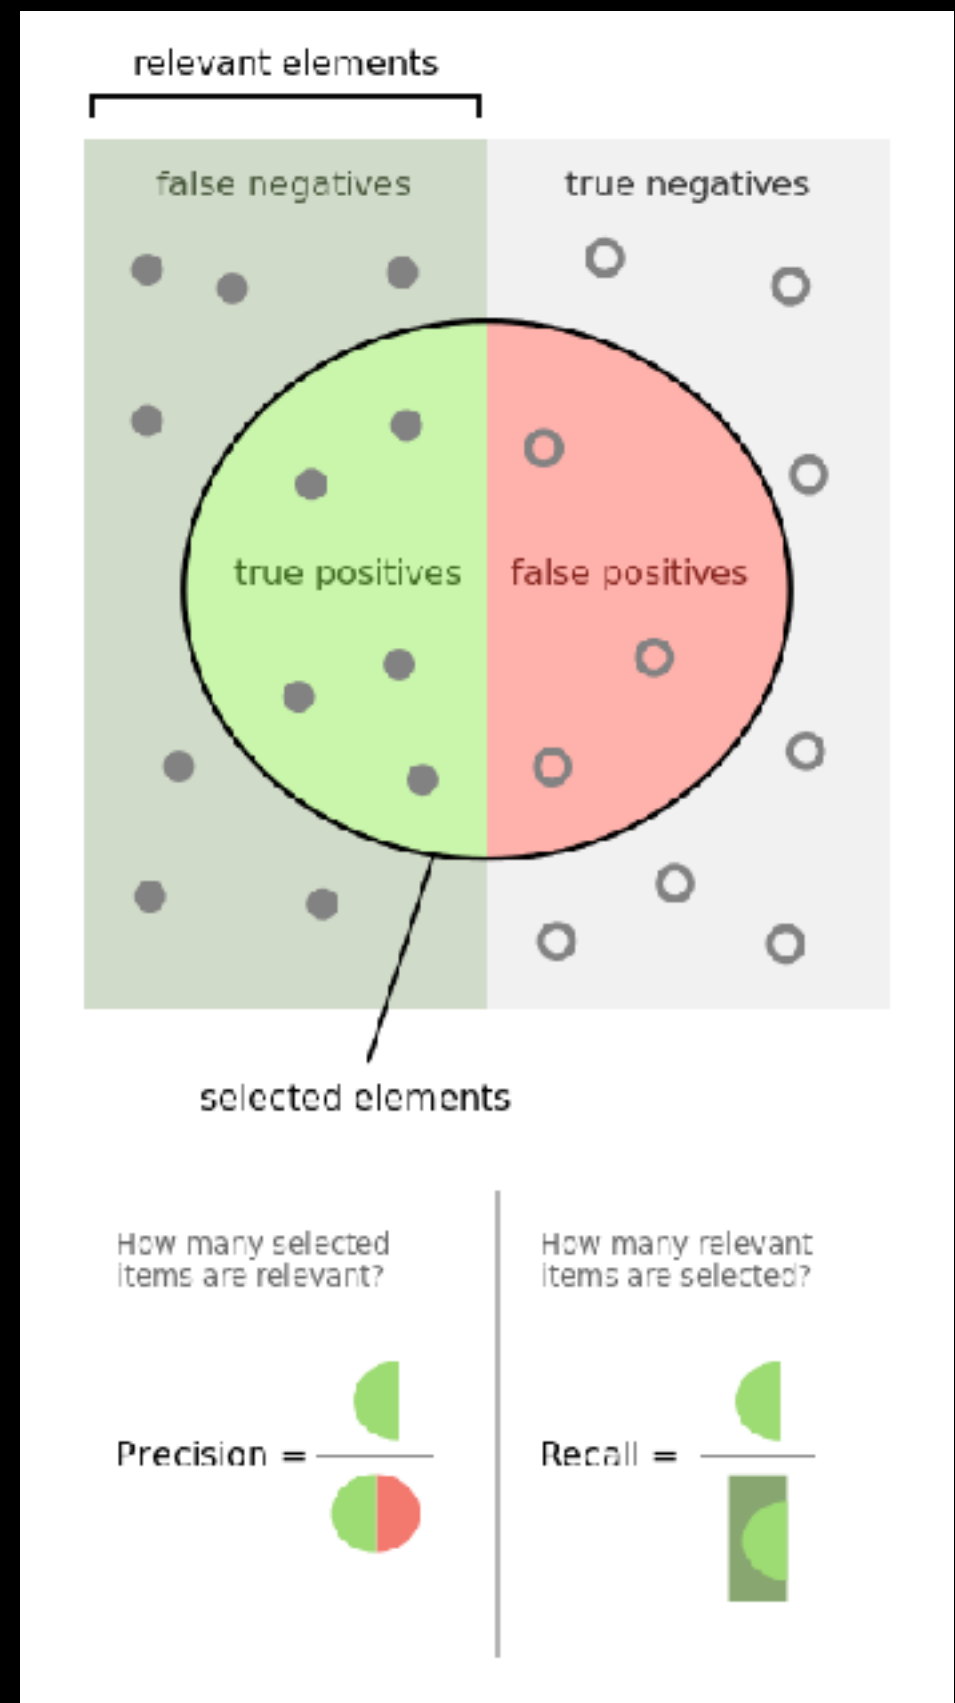
$$\text{Accuracy} = \frac{\text{TP} + \text{TN}}{\text{TP} + \text{TN} + \text{FP} + \text{FN}}$$

*TP = True positives*

*FP = False positives*

*TN = True negatives*

*FN = False negatives*



# Accuracy, Recall, Precision

$$\text{Precision} = \frac{\text{TP}}{\text{TP} + \text{FP}}$$

$$\text{Recall} = \frac{\text{TP}}{\text{TP} + \text{FN}}$$

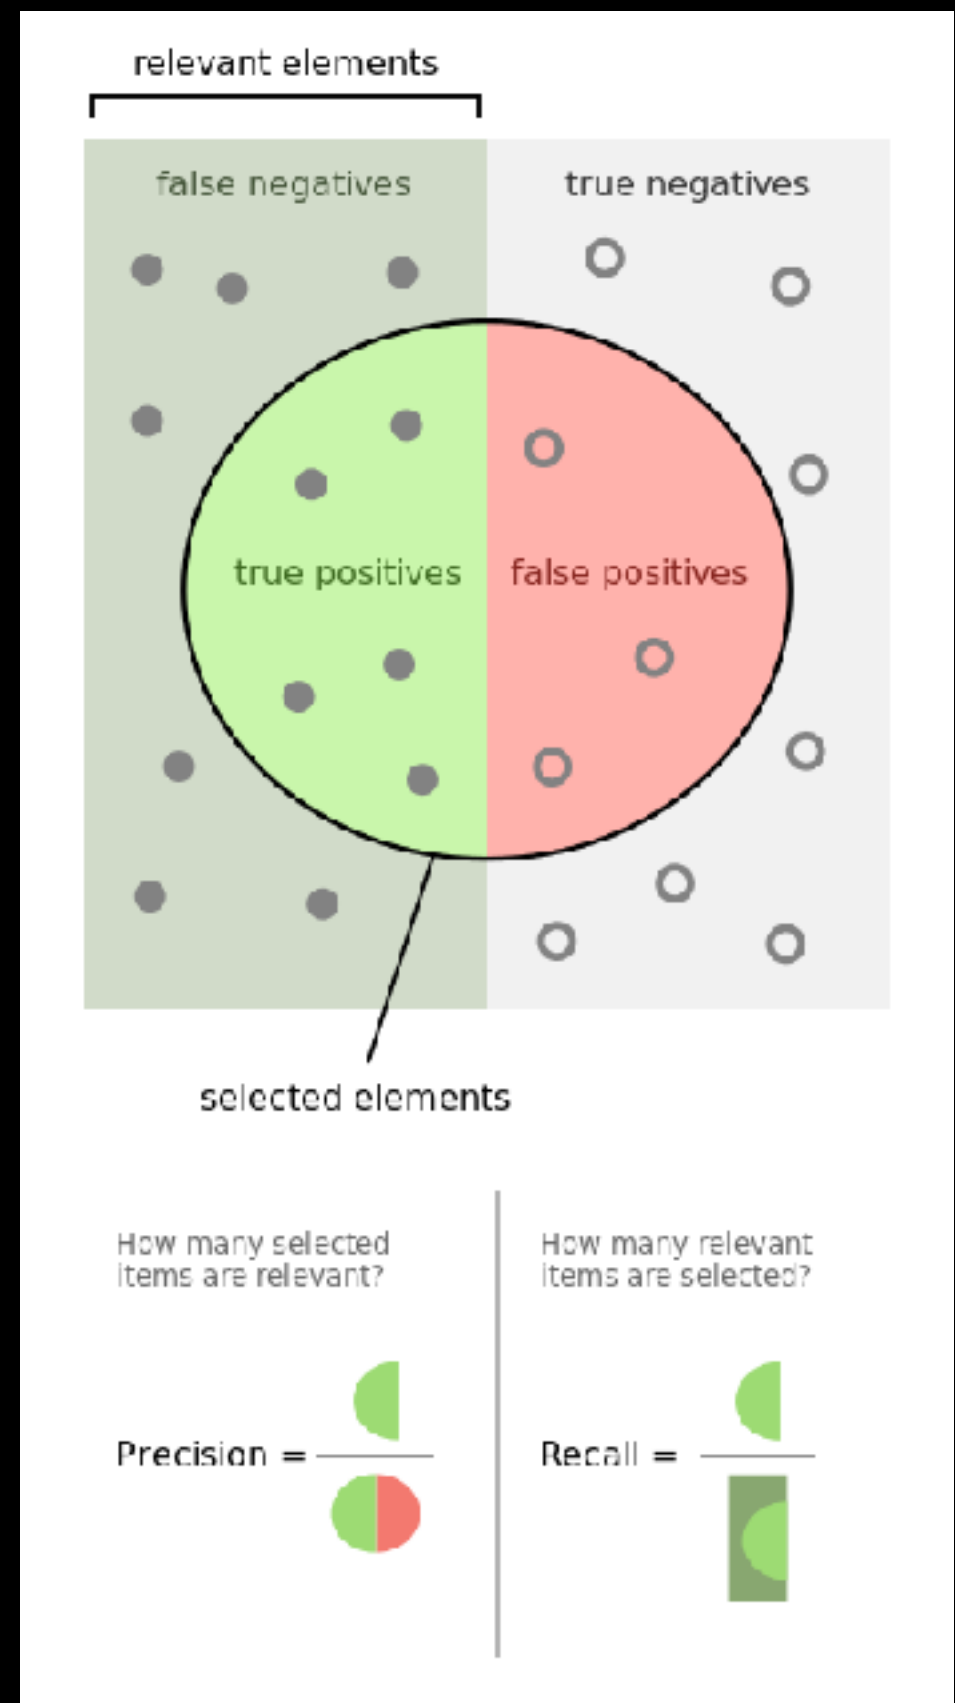
$$\text{Accuracy} = \frac{\text{TP} + \text{TN}}{\text{TP} + \text{TN} + \text{FP} + \text{FN}}$$

*TP = True positives*

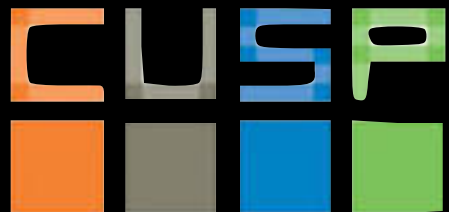
*FP = False positives*

*TN = True negatives*

*FN = False negatives*



# Data Preprocessing



XII categorical distances  
model diagnostics, OSMnx



# Full on whitening

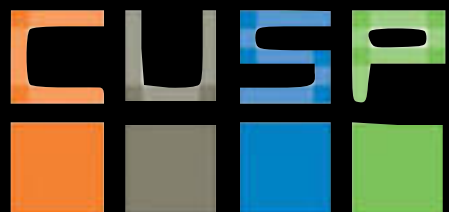
input data, PLUTO, manhattan (42,000x15)

axis 0 ↑

	Borough	Block	Lot	CD	CT2010	CB2010	SchoolDist	Council	ZipCode	FireComp	...	TaxMap	EDesignNum	APPBBL	APPDate	PLUTOMapID	Ver
0	MN	1545	52	108	138	4000	02	5	10028	E022	...	10515	None	0.000000e+00	None	1	
1	MN	723	7501	104	93	6000	02	3	10001	E003	...	10302	None	1.007230e+09	11/30/2006	1	
2	MN	1680	48	111	170	5000	04	8	10029	E091	...	10605	None	0.000000e+00	None	1	
3	MN	1385	32	108	130	2003	02	4	10021	E039	...	10508	None	0.000000e+00	None	1	
4	MN	1197	27	107	169	5000	03	6	10024	E074	...	10408	None	0.000000e+00	None	1	

axis 1 →

axis 0 : observations  
axis 1 : features



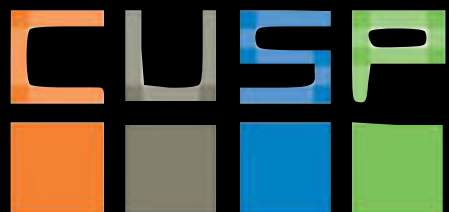
XII categorical distances  
model diagnostics, OSMnx

## Full on whitening

input data, PLUTO, manhattan (42,000x15)

$$\Sigma = \begin{bmatrix} E[(X_1 - \mu_1)(X_1 - \mu_1)] & E[(X_1 - \mu_1)(X_2 - \mu_2)] & \cdots & E[(X_1 - \mu_1)(X_n - \mu_n)] \\ E[(X_2 - \mu_2)(X_1 - \mu_1)] & E[(X_2 - \mu_2)(X_2 - \mu_2)] & \cdots & E[(X_2 - \mu_2)(X_n - \mu_n)] \\ \vdots & \vdots & \ddots & \vdots \\ E[(X_n - \mu_n)(X_1 - \mu_1)] & E[(X_n - \mu_n)(X_2 - \mu_2)] & \cdots & E[(X_n - \mu_n)(X_n - \mu_n)] \end{bmatrix}.$$

matrix of expected values of data



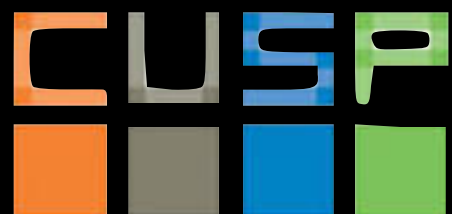
XII categorical distances  
model diagnostics, OSMnx

## Full on whitening

input data, PLUTO, manhattan (42,000x15)

$$\Sigma = \begin{bmatrix} E[(X_1 - \mu_1)(X_1 - \mu_1)] & 0 & \cdots & 0 \\ 0 & E[(X_2 - \mu_2)(X_2 - \mu_2)] & \cdots & 0 \\ \vdots & \vdots & \ddots & \vdots \\ 0 & 0 & & E[(X_n - \mu_n)(X_n - \mu_n)] \end{bmatrix}.$$

no covariance

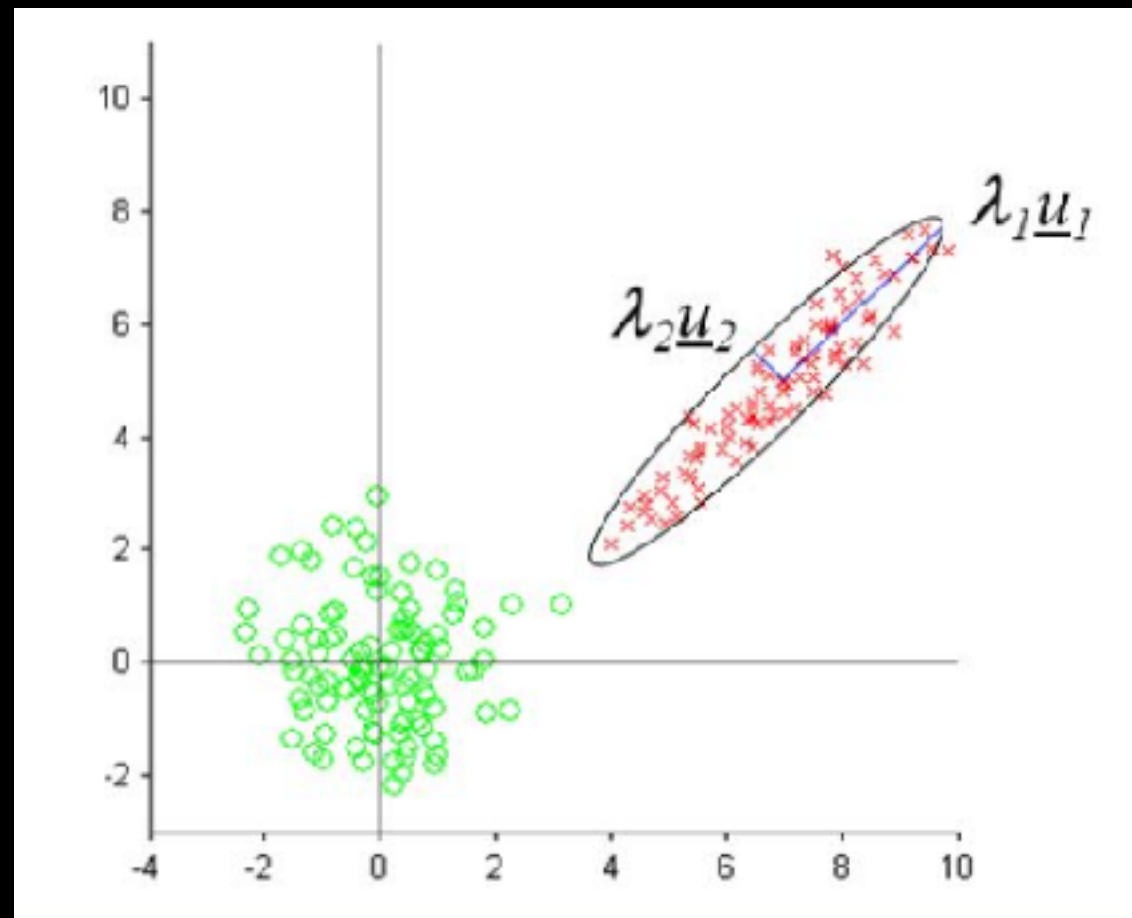


XII categorical distances  
model diagnostics, OSMnx

## Full on whitening

find the matrix  $W$  that diagonalized  $\Sigma$  covariance

$\Sigma W = \text{diagonal matrix}$

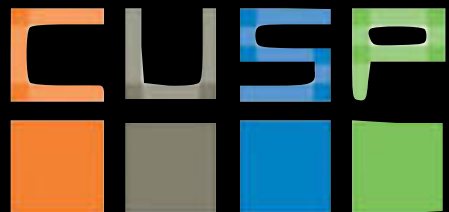


no covariance

## Full on whitening

find the matrix  $W$  that diagonalized  $\Sigma$

```
from zca import ZCA
import numpy as np
X = np.random.random((10000, 15)) # data array
trf = ZCA().fit(X)
X_whitened = trf.transform(X)
X_reconstructed = trf.inverse_transform(X_whitened)
assert(np.allclose(X, X_reconstructed)) # True
```



# Scaling - independent features

$X = \text{preprocessing.scale}(X)$

axis 0  $\uparrow$

	Borough	Block	Lot	CD	CT2010	CB2010	SchoolDist	Council	ZipCode	FireComp	...	TaxMap	EDesignNum	APPBBL	APPDate	PLUTOMapID	Ver
0	MN	1545	52	108	138	4000	02	5	10028	E022	...	10515	None	0.000000e+00	None	1	
1	MN	723	7501	104	93	6000	02	3	10001	E003	...	10302	None	1.007230e+09	11/30/2006	1	
2	MN	1680	48	111	170	5000	04	8	10029	E091	...	10605	None	0.000000e+00	None	1	
3	MN	1385	32	108	130	2003	02	4	10021	E039	...	10508	None	0.000000e+00	None	1	
4	MN	1197	27	107	169	5000	03	6	10024	E074	...	10408	None	0.000000e+00	None	1	

axis 0  $\longrightarrow$  mean = 0, stdev = 1

```
X = preprocessing.scale(X, axis=0)
```

Last executed 2018-12-12 09:35:39 in 46ms

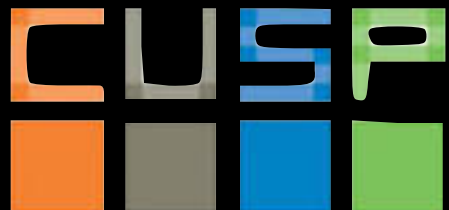
```
X.mean(axis=0)
```

Last executed 2018-12-12 09:35:40 in 13ms

```
array([ 3.85590369e-16, -6.93196168e-17, -5.90549813e-16, -5.95882091e-16,
       -8.49165306e-16, -1.57568821e-15, -8.00508267e-16,  5.55890004e-16,
       -5.16564452e-16,  1.09378357e-15,  3.46598084e-16,  2.31954102e-16,
        2.78611537e-16, -2.51283611e-16,  8.66495210e-18,  3.03939858e-16,
       -3.66594127e-17, -9.27149875e-16, -6.39873386e-16,  2.93275302e-17,
        9.19817992e-17,  6.33208038e-18, -1.99960433e-17,  9.55144336e-16,
       -2.20623011e-16,  6.93196168e-17, -9.46479383e-17,  2.26621824e-16,
        6.93196168e-17,  2.32953905e-16])
```

HW11

XII categorical distances  
model diagnostics, OSMnx



# Scaling - independent features

```
X = preprocessing.scale(X)
```

	Borough	Block	Lot	CD	CT2010	CB2010	SchoolDist	Council	ZipCode	FireComp	...	TaxMap	EDesignNum	APPBBL	APPDate	PLUTOMapID	Ver
0	MN	1545	52	108	138	4000	02	5	10028	E022	...	10515	None	0.000000e+00	None	1	
1	MN	723	7501	104	93	6000	02	3	10001	E003	...	10302	None	1.007230e+09	11/30/2006	1	
2	MN	1680	48	111	170	5000	04	8	10029	E091	...	10605	None	0.000000e+00	None	1	
3	MN	1385	32	108	130	2003	02	4	10021	E039	...	10508	None	0.000000e+00	None	1	
4	MN	1197	27	107	169	5000	03	6	10024	E074	...	10408	None	0.000000e+00	None	1	

axis 0  $\longrightarrow$

axis 0  $\longrightarrow$  mean = 0, stdev = 1

```
X = preprocessing.scale(X, axis=0)
```

Last executed 2018-12-12 09:35:39 in 46ms

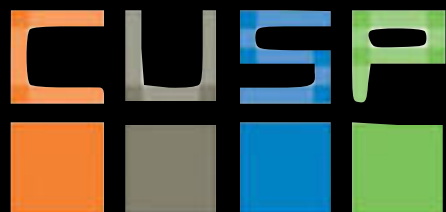
```
X.std(axis=0)
```

Last executed 2018-12-12 09:36:28 in 19ms

```
array([1., 1., 1., 1., 1., 1., 1., 1., 1., 1., 1., 1., 1., 1., 1.,  
       1., 1., 1., 1., 1., 1., 1., 1., 1., 1., 1., 1.])
```

# HW11

## XII categorical distances model diagnostics, OSMnx



# Scaling - feature vectors: *e.g. time series*

`X = preprocessing.scale(X, axis=1)`

	Borough	Block	Lot	CD	CT2010	CB2010	SchoolDist	Council	ZipCode	FireComp	...	TaxMap	EDesignNum	APPBBL	APPDate	PLUTOMapID	Ver
0	MN	1545	52	108	138	4000	02	5	10028	E022	...	10515	None	0.000000e+00	None	1	
1	MN	723	7501	104	93	6000	02	3	10001	E003	...	10302	None	1.007230e+09	11/30/2006	1	
2	MN	1680	48	111	170	5000	04	8	10029	E091	...	10605	None	0.000000e+00	None	1	
3	MN	1385	32	108	130	2003	02	4	10021	E039	...	10508	None	0.000000e+00	None	1	
4	MN	1197	27	107	169	5000	03	6	10024	E074	...	10408	None	0.000000e+00	None	1	

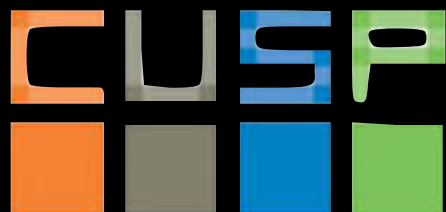
axis 1  $\longrightarrow$  mean = 0, stdev = 1

**Build one that uses as input features the following engineered features :**

- the time series mean divided by the mean of all time series for that station
- the time series standard deviation by the standard deviation of all time series for that station
- the slope and intercept of a line fit to the standardized time series

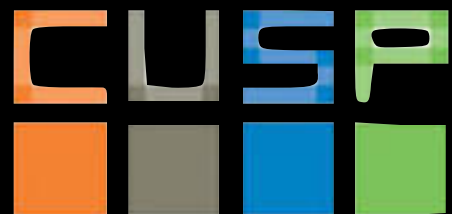
```
(time_series - time_series.mean())/time_series.std()
```

you will have to remove time series containing NaN because the random forest sklearn implementation does not work with NaNs. An easy way to do that is to remove all time series whose standard deviation is NaN





super important missing topic:  
**pruning!**  
when is my tree overfitting?



## is your code optimized:

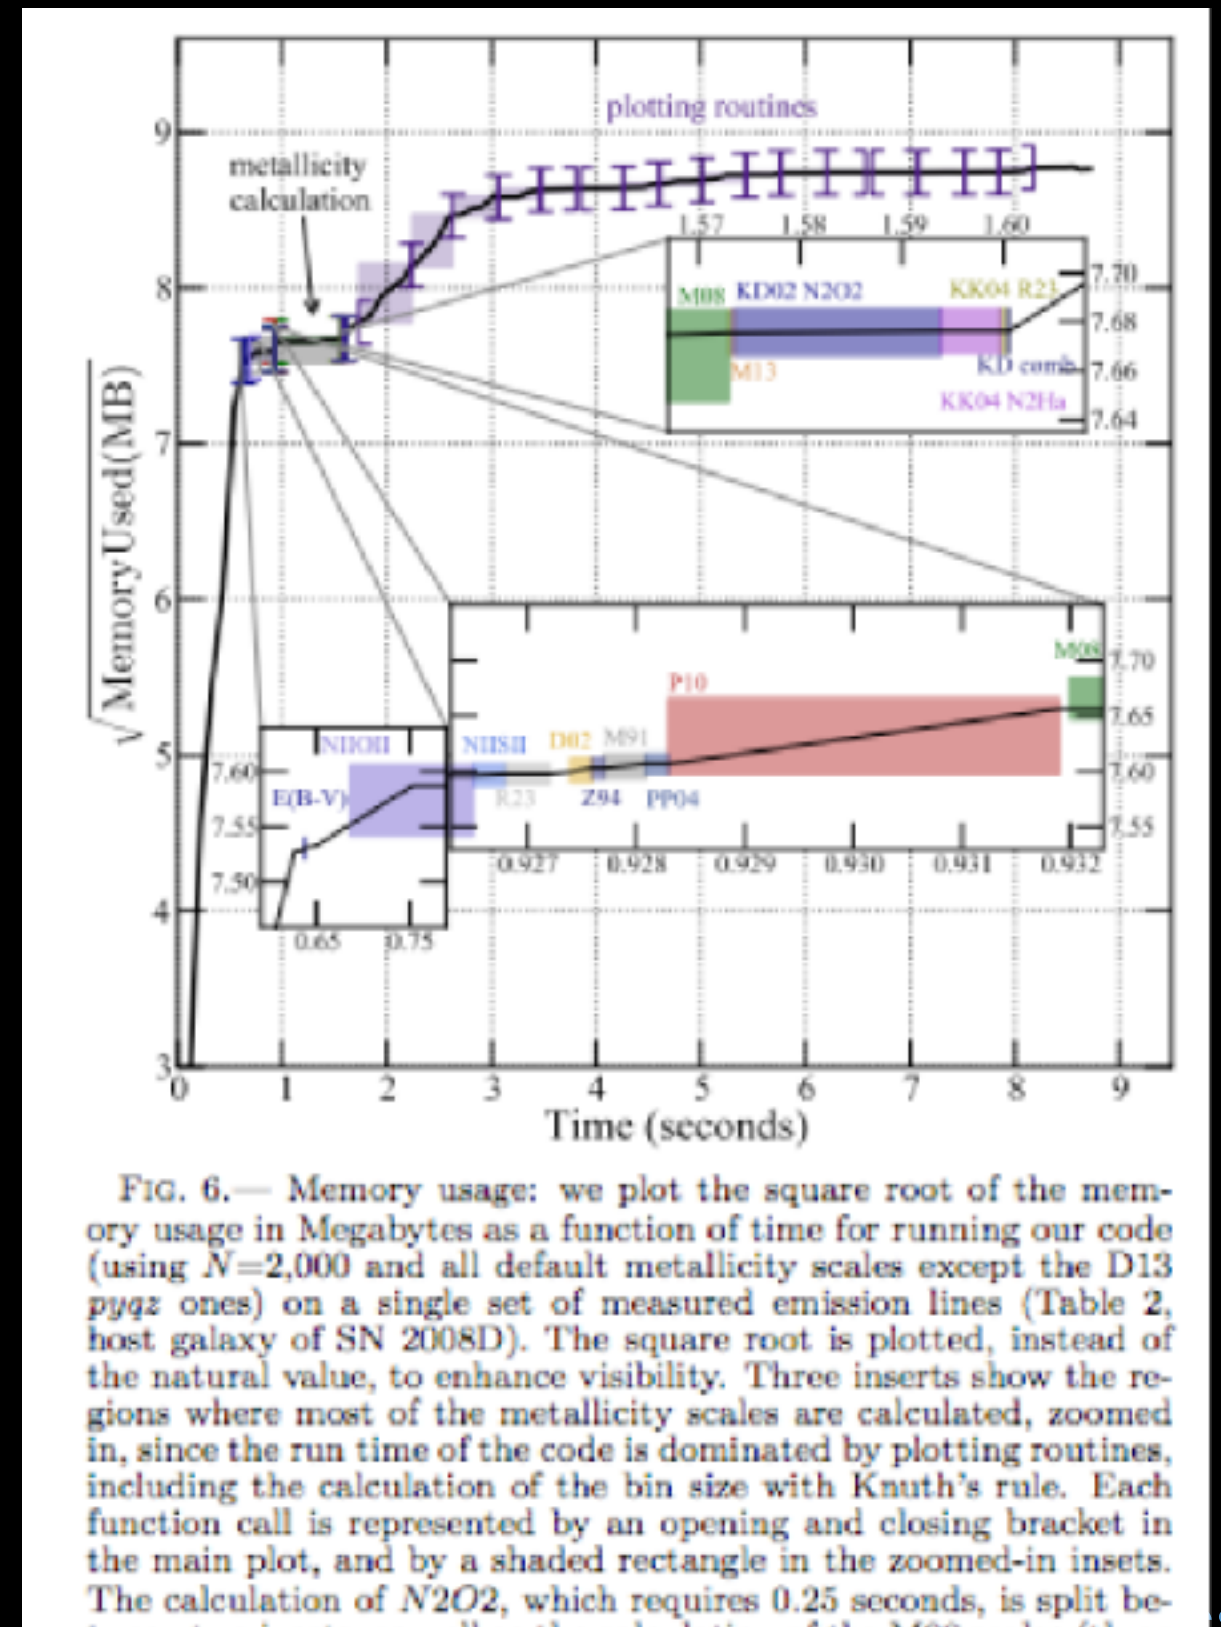
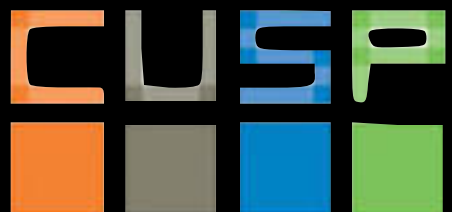
check CPU AND MEMORY usage

vectorize (slice and avoid for loops)

avoid storing information you do not need in memory

use local variables

remove all redundant calculations from inside loops

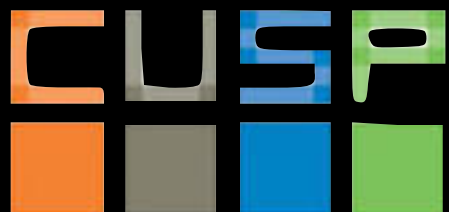


## Reading:

*An excellent use of viz for data exploration  
and transition to inferential analysis*

<https://blog.data.gov.sg/how-we-caught-the-circle-line-rogue-train-with-data-79405c86ab6a#.iz1r655xo>

Lee Shangqian, Daniel Sim & Clarence Ng



XII categorical distances  
model diagnostics, OSMnx

## Distance measures for clustering:

[http://sfb649.wiwi.hu-berlin.de/fedc\\_homepage/xplore/tutorials/mvahtmlnode79.html](http://sfb649.wiwi.hu-berlin.de/fedc_homepage/xplore/tutorials/mvahtmlnode79.html)

## Decision trees:

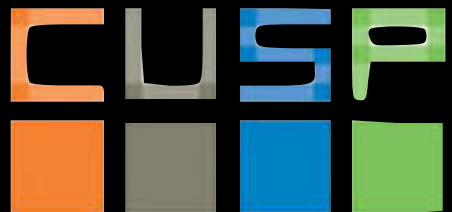
<http://what-when-how.com/artificial-intelligence/decision-tree-applications-for-data-modelling-artificial-intelligence/>

<https://www.ncbi.nlm.nih.gov/pmc/articles/PMC4466856/>

<https://www.ncbi.nlm.nih.gov/pmc/articles/PMC4380222/>

## Efficient python coding:

<https://wiki.python.org/moin/PythonSpeed/PerformanceTips>



XII categorical distances  
model diagnostics, OSMnx